

Nimbus AI Suite -- Product Handbook v2.1

Overview

The Nimbus AI Suite is a document-automation platform developed by Acme Corp's Software division. It combines hierarchical retrieval-augmented generation with workflow automation to help enterprise teams extract, summarize, and act on information buried in large document sets.

This handbook covers version 2.1, released in October 2025.

System Requirements

The Nimbus AI Suite desktop client requires:

- 8 GB RAM minimum (16 GB recommended)
- Windows 10/11, macOS 12+, or a modern Linux distribution
- Python 3.9 or later for the optional command-line integration
- A broadband internet connection for cloud-synced workspaces

The Nimbus AI Suite server component additionally requires at least 4 CPU cores and 50 GB of available disk space for document storage.

Pricing Tiers

Nimbus AI Suite is offered in three pricing tiers, billed monthly:

Starter: \$29 per user per month -- up to 500 documents

Pro: \$99 per user per month -- up to 10,000 documents, priority support

Enterprise: custom pricing -- unlimited documents, dedicated support, and a guaranteed service-level agreement

Annual billing is available on the Pro and Enterprise tiers at a 15% discount.

Key Features

Nimbus AI Suite's core features include:

- Hierarchical chunking: large 'parent' sections for context, small 'child' passages for precise retrieval
- Multi-document support: ingest and query across an entire library of PDFs at once, with per-document citation tracking
- Cross-encoder re-ranking for higher-precision search results
- Automatic citation generation, linking every answer back to the exact source page it came from

Support

Nimbus AI Suite customers can reach support at support@nimbusai.example.com.

Service-level agreement (SLA) response times by tier:

Starter: best-effort, no guaranteed response time

Pro: 48-hour response time

Enterprise: 24-hour response time, with a dedicated account manager

Status updates and planned maintenance windows are posted to the Nimbus AI status page.

Changelog

Version 2.1 (October 2025): added a dark mode theme and improved PDF table extraction accuracy.

Version 2.0 (July 2025): added native integration with the Groq inference API for faster answer generation; this was the release referenced in Acme's Q3 2025 financial report as the quarter's flagship product launch.

Version 1.0 (January 2024): initial public release, supporting single-document ingestion only.